

# eBPF Misbehavior Detection: Fuzzing with a Specification-Based Oracle


Tao Lyu Kumar Kartikeya Dwivedi

Thomas Bourgeat Mathias Payer Meng Xu\* Sanidhya Kashyap




# Kernel Extensions Are Essential

- Extending kernels at runtime with safety guarantees
- Widely deployed for observability, security, networking, etc.



# Two Tiny eBPF Programs, Two Big Verifier Issues

- Security issues




```
...
r1 = r10
if cond goto pc+2
*(u64*)(r1-120) = 0
goto pc+1
...
```

```
tlyu@syzkaller:~$ ./exp-get-root
[+] Created BPF maps and confirmed BPF is enabled on this kernel.
[+] OOB read succeeded! The kernel is vulnerable
[+] Now we get every rights. See you on the other side!
# id
uid=0(root) gid=0(root) groups=0(root),1000(tlyu) context=system_u:system_r:kernel_t:s0
# whoami
root
#
```




# Two Tiny eBPF Programs, Two Big Verifier Issues

- Usability issues



# So... What's Wrong with the eBPF Verifier?



over-approximated state

incorrect checks


mis-implementations

# Verifier Internals & Root Causes

verifier = path enumeration + abstract interpretation + safety checks


# Verifier Internals & Root Causes

verifier = path enumeration + abstract interpretation + safety checks



# Verifier Internals & Root Causes

verifier = path enumeration + abstract interpretation + safety checks




## abstract state

r10: stack pointer, r4: a pointer  
r2: [100, 255], r3: [200,255]

# Verifier Internals & Root Causes

verifier = path enumeration + abstract interpretation + safety checks



## abstract state

$r10$ : stack pointer,  $r4$ : a pointer  
 $r2$ :  $[100, 255]$ ,  $r3$ :  $[0, 2^{64}-1]$


$r3: [0,2] \text{ vs } [0, 2^{64}-1]$

E.g., Out of bound access (OOB) if use  
 $r3$  as an index for 20 byte array

over-approximated state

# Verifier Internals & Root Causes

verifier = path enumeration + abstract interpretation + safety checks




## abstract state

r10: stack pointer, r4: a pointer  
r2: [100, 255], r3: [8,  $2^{64}-1$ ]

over-approximated state

# Verifier Internals & Root Causes

verifier = path enumeration + abstract interpretation + safety checks



## abstract state


r10: stack pointer, r4: a pointer  
r2: [100, 255], r3: [8,  $2^{64}-1$ ]

over-approximated state

incorrect checks

# Verifier Internals & Root Causes

verifier = path enumeration + abstract interpretation + safety checks



## abstract state


$r10$ : stack pointer,  $r4$ : a pointer  
 $r2$ :  $[100, 255]$ ,  $r3$ :  $[0, 8]$

over-approximated state

incorrect checks

# Verifier Internals & Root Causes

verifier = path enumeration + abstract interpretation + safety checks



## abstract state

$r10$ : stack pointer,  $r4$ : a pointer  
 $r2$ : [100, 255],  $r3$ : [0, 8]

## safety checks


within-bound mem access

over-approximated state

incorrect checks

# Verifier Internals & Root Causes

verifier = path enumeration + abstract interpretation + safety checks



## abstract state

r10: stack pointer, r4: a pointer  
r2: [100, 255], r3: [0, 8]

## safety checks

within-bound mem access

## path pruning

if  $s2 \subseteq s1$ , skip last code block in **2**

over-approximated state

incorrect checks

mis-implementations

# Remedies to the Verifier's Issues

	Fuzzing with heuristics-based checker (e.g., KASAN)
Over-approximated state	✗
Incorrect safety checks	✗
Mis-implementations	✗✓

# Remedies to the Verifier's Issues

	Fuzzing with heuristics-based checker (e.g., KASAN)	Partial verification (i.e., Agni)
Over-approximated state	✗	✓
Incorrect safety checks	✗	✗
Mis-implementations	✗	✓

**They all have limited capability to address these issues!**

# Our Goal: Heuristic → Systematic

## CORE QUESTION

What is the **correctness** of the eBPF verifier?




## SPECIFICATION

=

instruction semantics + safety constraints

# Veritas Roadmap



# What the Specification Looks Like?

- Specification = per-instruction semantics + safety constraints

```
function insn(s: State, insn: Instruction) : State
```

```
// safety constraints as preconditions  
requires ...
```

```
{
```

```
// semantics as a function body
```

```
}
```

# What the Specification Looks Like?

- Specification = per-instruction semantics + safety constraints

```
function neg32(s: State, insn: Instruction) : State  
  
{  
  
}  
}
```

# What the Specification Looks Like?

- Specification = per-instruction semantics + **safety constraints**

```
function neg32(s: State, insn: Instruction) : State  
  
    // safety constraints as preconditions  
    requires insn.dst != R10 && ...  
  
{  
  
}
```

# What the Specification Looks Like?

- Specification = per-instruction **semantics** + safety constraints

```
function neg32(s: State, insn: Instruction) : State  
  
// safety constraints as preconditions  
requires insn.dst != R10 && ...  
{  
    // semantics as a function body  
    var new_val := bvnot32(arith_val(s, insn.dst));  
}  
}
```

# Augmenting Semantics with Dynamic Types

- Prevents value-semantics mismatches
  - E.g., using an integer instead of a pointer to access memory

# Augmenting Semantics with Dynamic Types

- Prevents value-semantics mismatches
  - E.g., using an integer instead of a pointer to access memory
- Each register and memory byte is associated with a dynamic type

```
datatype TYPEDVAL =
| Uninit
| Scalar(...)
| PtrType(...)
| PtrOrNullType(...)
```

# Augmenting Semantics with Dynamic Types

- Prevents value-semantics mismatches
  - E.g., using an integer instead of a pointer to access memory
- Each register and memory byte is associated with a dynamic type
- Each instruction corresponds to a type rule = 
$$\frac{\text{required types of src operands}}{\text{produced type of dst operands}}$$

```
datatype TYPEDVAL =
    | Uninit
    | Scalar(...)
    | PtrType(...)
    | PtrOrNullType(...)
```


# Augmenting Semantics with Dynamic Types

- Specification = (**augmented**) per-instruction semantics + safety constraints

```
function neg32(s: State, insn: Instruction) : State
  requires get_reg_tv(s, insn.dst) != Uninit
  // safety constraints as preconditions
  requires insn.dst != R10
  {
    // semantics as a function body
    var new_val := bvnot32(arith_val(s, insn.dst));
    new_state(s, insn.dst, Scalar(Normal, new_val))
  }
```

# Mid-talk Recap

- Specification = per-instruction semantics + safety constraints



# Where the Safety Constraints Come From?


## CIA triad security model for the OS kernel

### Confidentiality

$$s_1 \underset{\approx}{\underset{L}{\approx}} s_2$$

|  
insn  
↓  
 $s'_1 \underset{\approx}{\underset{L}{\approx}} s'_2$

### Integrity



### Availability

- Termination
- No resource exhaustion
- No memory crashes

high-security data  $\not\Rightarrow$  low    Non-writable data is unmodified


# Specification: Safety Constraints

## Confidentiality

$$s_1 \xrightarrow[L]{\quad} s_2$$

|  
insn  
↓  
 $s'_1 \xrightarrow[L]{\quad} s'_2$

## Integrity



## Availability

- Termination
- No resource exhaustion
- No memory crashes

Control flow safety

VM invariant

Memory safety

VM data safety

Resource safety


# Specification: Safety Constraints

## Confidentiality

$$s_1 \xrightarrow[L]{\quad} s_2$$

|  
insn  
↓  
 $s'_1 \xrightarrow[L]{\quad} s'_2$

## Integrity



## Availability

- Termination
- No resource exhaustion
- No memory crashes

Control flow safety

Memory safety

Resource safety

**Please check out the detailed safety constraints in our paper**

VM invariant

VM data safety

# Specification Soundness and Completeness

Formal

Confidentiality

Soundness ✓

Integrity

Soundness ✓


Availability

Empirical

- No false alarms in the fuzzing campaign
- eBPF selftests: pass, or discrepancy with a known cause

**Specification ➔ Improving the verifier's correctness**

# Specification as a Fuzzing Oracle



# Translating eBPF Bytecode to Dafny

eBPF assembly code	Dafny code
r2 = *(u64*) (r1 + 0)	s0 := init_vm_state(cfg)
if r2 == 0 goto L1	s1 := mem_load(s0, MEMLD(R2, R1, 0, )) s2 := jeq(s1, CONDJMPIMM(R2, 0, ))
...	if (!s2.jmp_res) {
call	...
L1:	s4 := map_lookup_elem(s3'); exit(s4);
exit	} else {exit(s2);}

# Evaluation

- Uncovered 13 new bugs
  - 3 security issues + 10 usability issues
- Comparison
  - SpecCheck detects all existing 14 bugs
  - Baselines cannot detect bugs found by SpecCheck
- Throughput
  - 23-25 tests / second
  - Trade throughput for better bug-detection capability

# Conclusion

- The Linux eBPF verifier is critical but vulnerable
- Fuzzing with a specification-based oracle
  - Specification = semantics + safety constraints
  - eBPF bytecodes  $\xrightarrow{\text{spec}}$  Dafny programs  $\xrightarrow{\text{Dafny verifier}}$  oracles
- Future use cases
  - Proof-carrying code (PCC); verifying user-defined properties
- Artifact: <https://github.com/rs3lab/veritas>

Thank you!