

Tolerate It if You Cannot Reduce It: Handling Latency in Tiered Memory

Musa Unal, Vishal Gupta, Yueyang Pan, Yujie Ren, Sanidhya Kashyap



Tiered memory

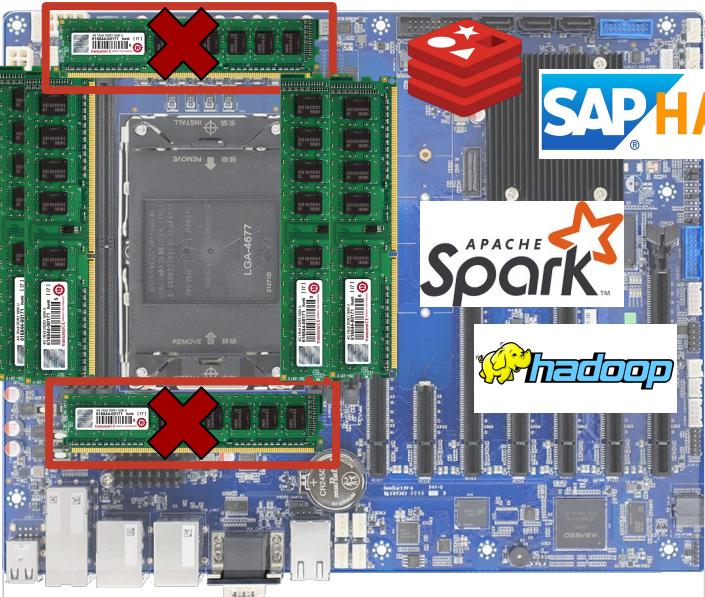
Why ?

Applications demand for memory.

We want to expand main memory.

We can not add more DRAMs.

How ?

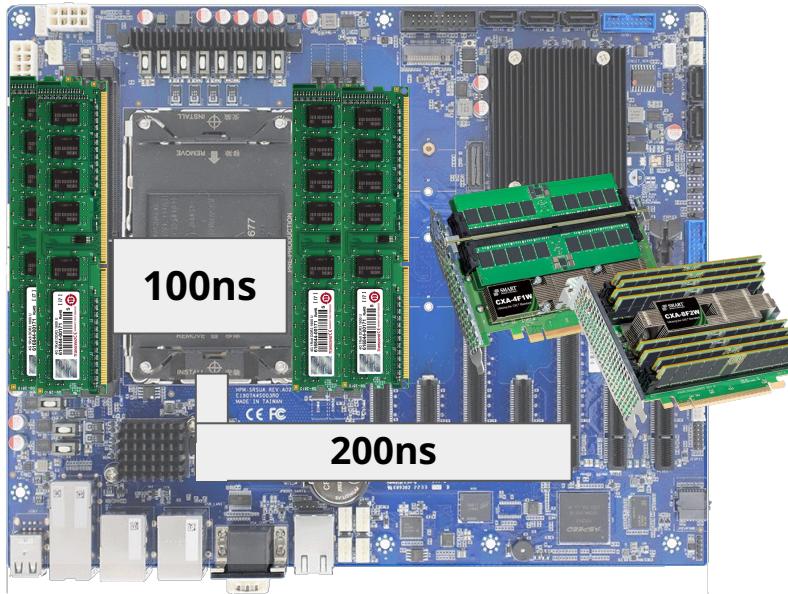


Memory

DRAM

CXL

CXL expands main memory



CXL Attached Memory

Memory

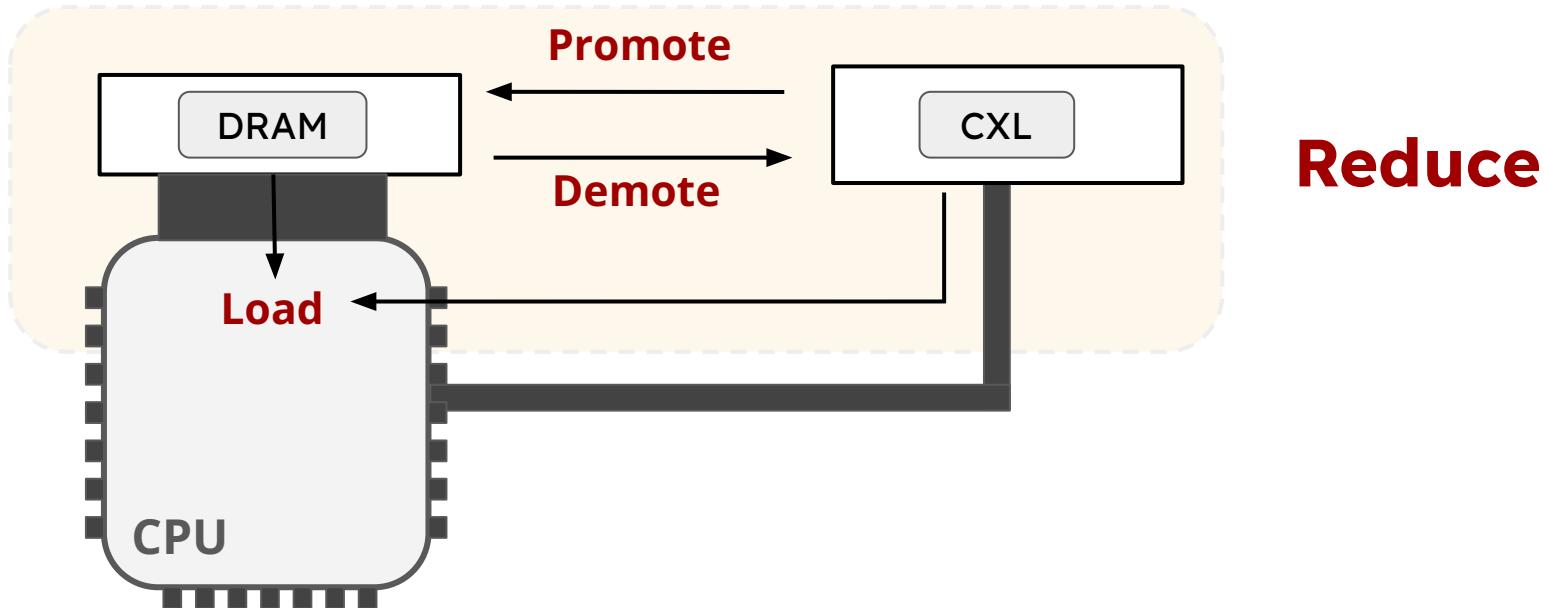
DRAM

CXL

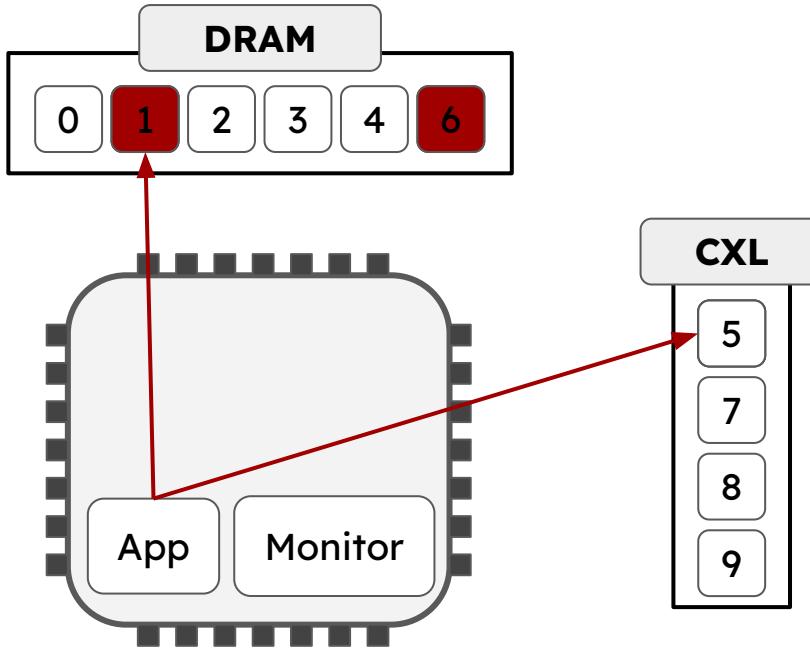
- Protocol for fast memory access over PCIe
- Latency is 2x higher (or higher with different topologies)

How can we handle the latency in tiered memory?

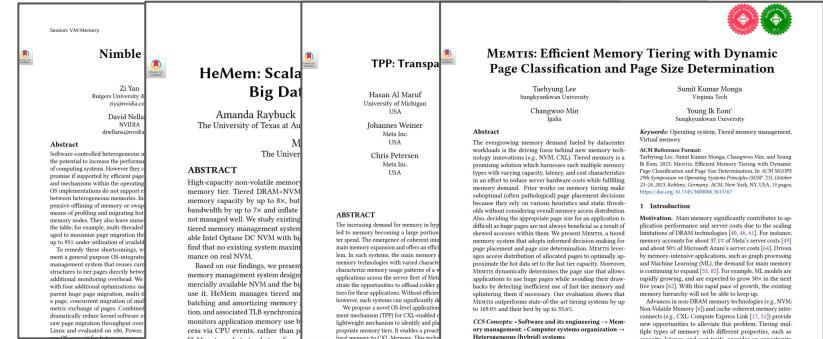
Conventional approach to handle latency



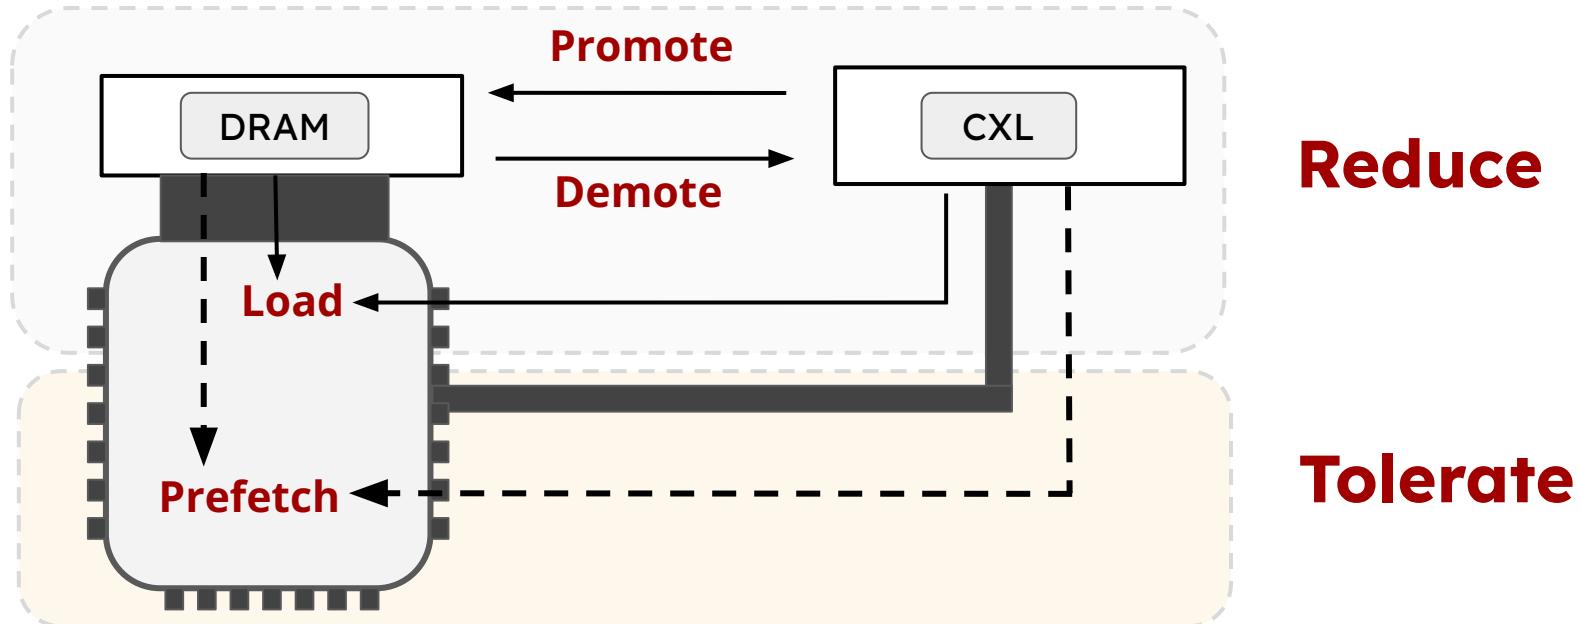
Page migration can reduce the latency



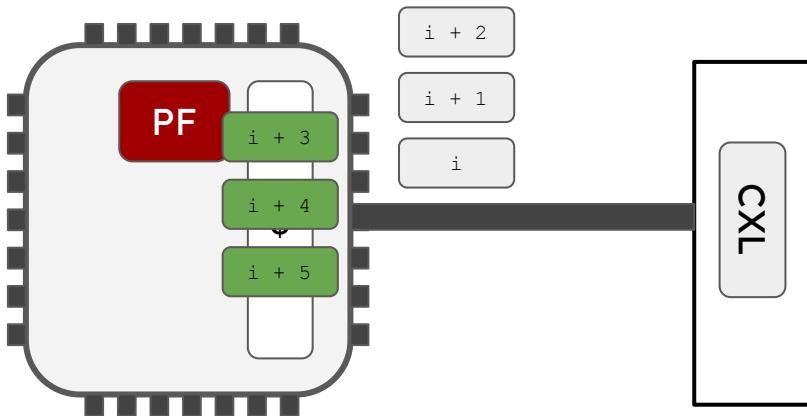
- 1 Tiering systems monitor pages hotness information.
 - 2 Promote hot pages
Demote cold pages
 - 3 Increase **locality**, and
reduce the latency



Missing aspect to handle latency



Hardware prefetchers can tolerate latency

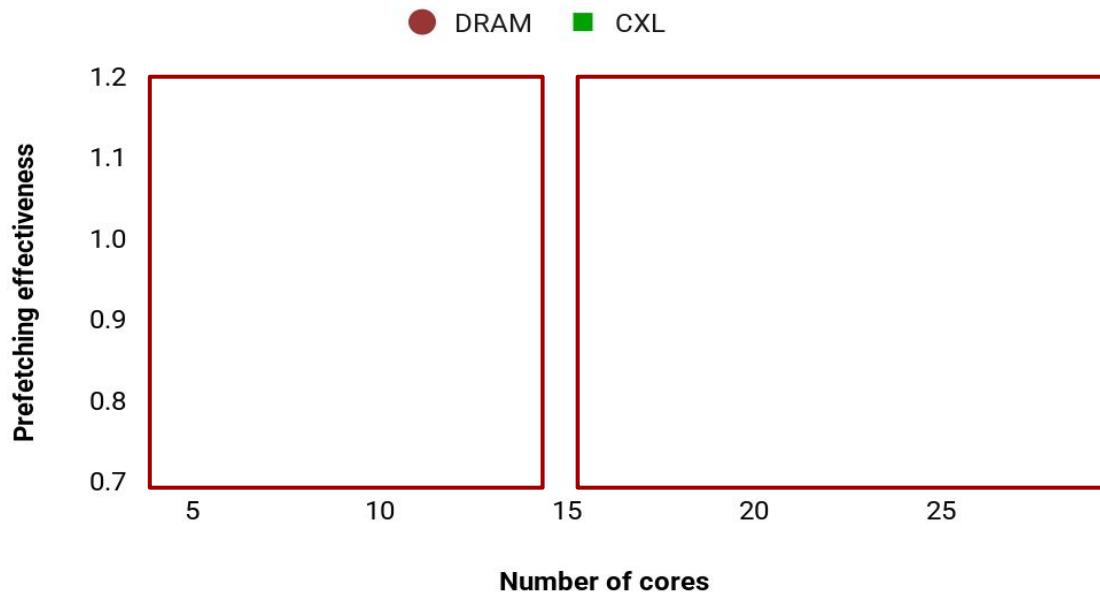


Sequential Pointer Chase

Stride Size	CXL Lat / DRAM Lat
4096	~2.01
128	~1.79
64	~1.55

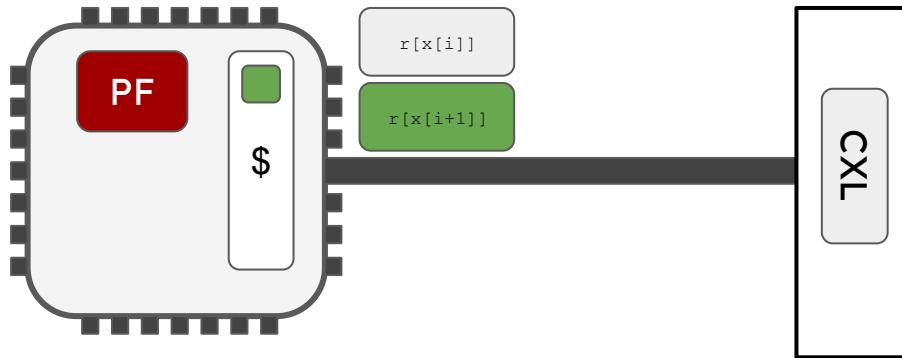
Hardware prefetchers can help to hide CXL's latency.

Hardware prefetchers can be harmful



Prefetching can degrade the performance under high loads.

Software prefetching can help as well



Timeliness is important!

Early prefetches can be evicted.

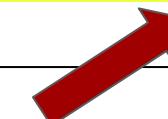
Late prefetches are not helpful.

```
for i in range(1M):  
    e = records[x[i]]  
    work(e)
```

```
for i in range(1M):  
    e = records[x[i]]  
    prefetch(records[x[i+1]])  
    work(e)
```

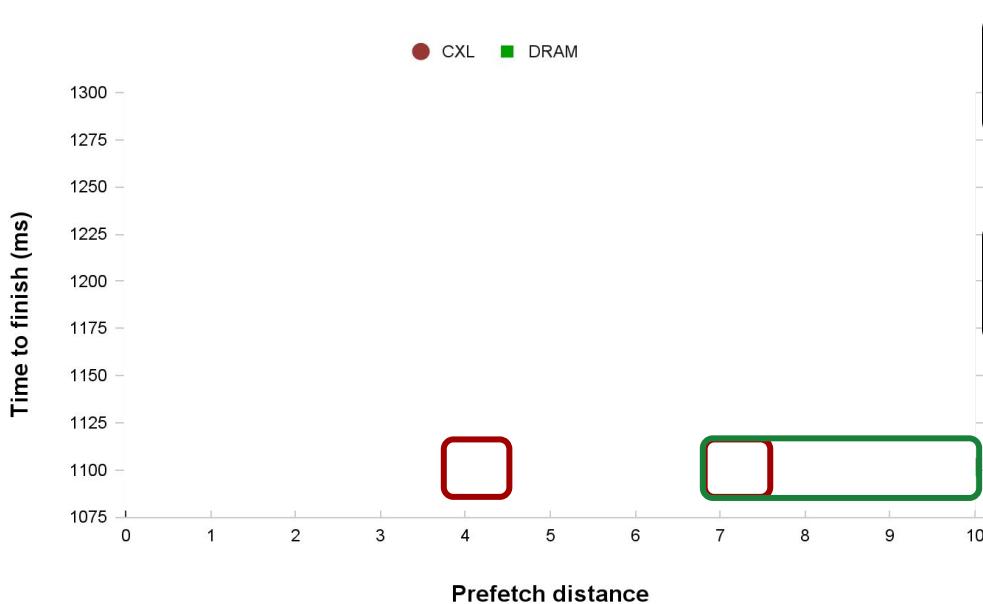
```
for i in range(1M):  
    e = records[x[i]]  
    prefetch(records[x[i+k]])  
    work(e)
```

Prefetch distance



Prefetch distance matters

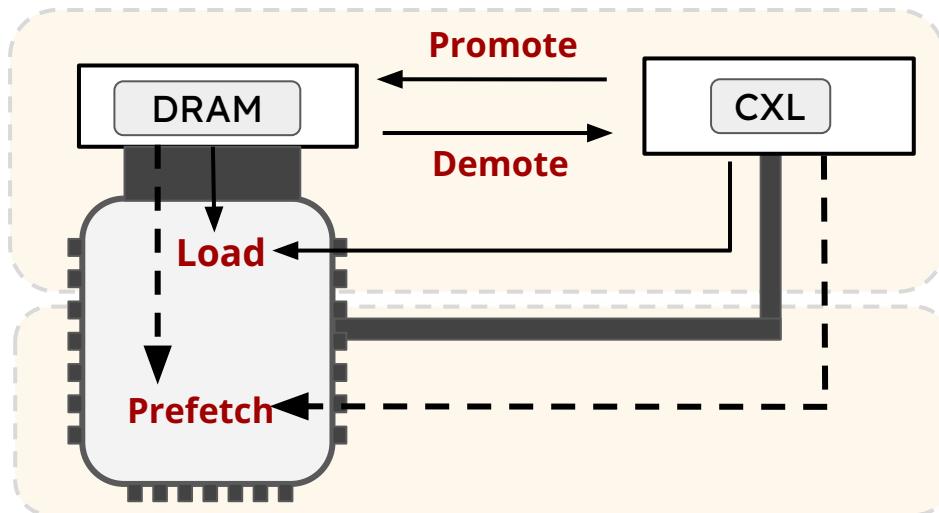
Scan microbenchmark



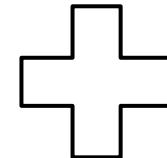
Prefetch distance should differ across tiers

Prefetching can effectively **hide** the latency

Linden: Combining reduce and tolerate



Reduce



→ Hotness

Tolerate

→ Prefetchability

Prefetchability as a tolerate metric

Can we identify memory regions that will get benefit from prefetching ?

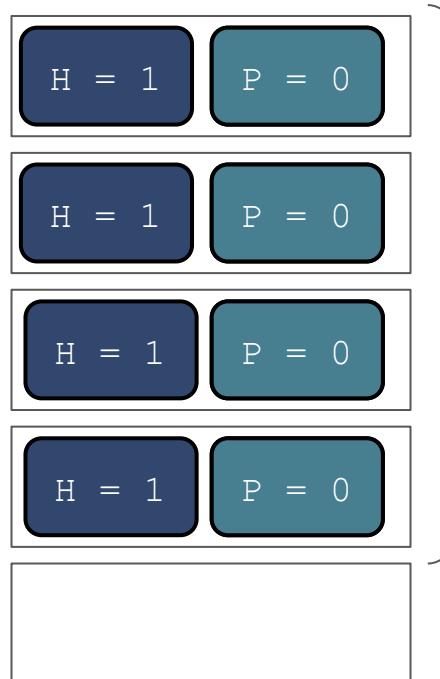
Prefetchability:

Percentage of the loads to particular region that can be hidden by prefetching.

$$\frac{\text{Number of loads hidden with prefetching}}{\text{Issued loads}}$$

Thread	Region	Access Pattern	Prefetchability Ratio
T_1	arr	Sequential	1
T_2	arr_2	Random	0

Combining hotness and prefetchability



```
// rand access
arr = [...]
for j in range(10):
    for i in range(1000):
        sum += arr[rand]
```

H=0-1

Hotness

P=0-1

Prefetchability

Memory Pages

Linden's objectives and workflow

Prefetching can hide CXL's latency.

Under high contentions prefetching can be harmful

Prefetch distances matter in different tiers

Compiler

Identify prefetchable access patterns

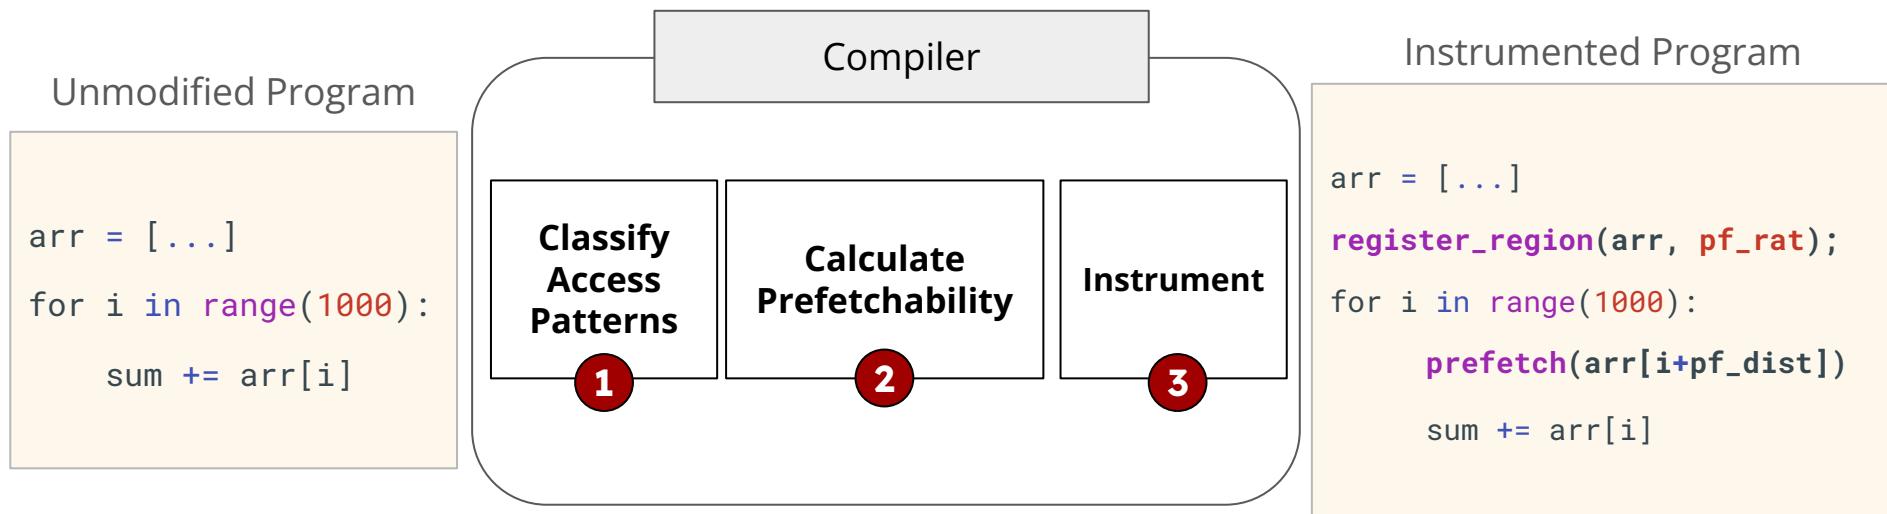
Insert software prefetch hints.

Runtime

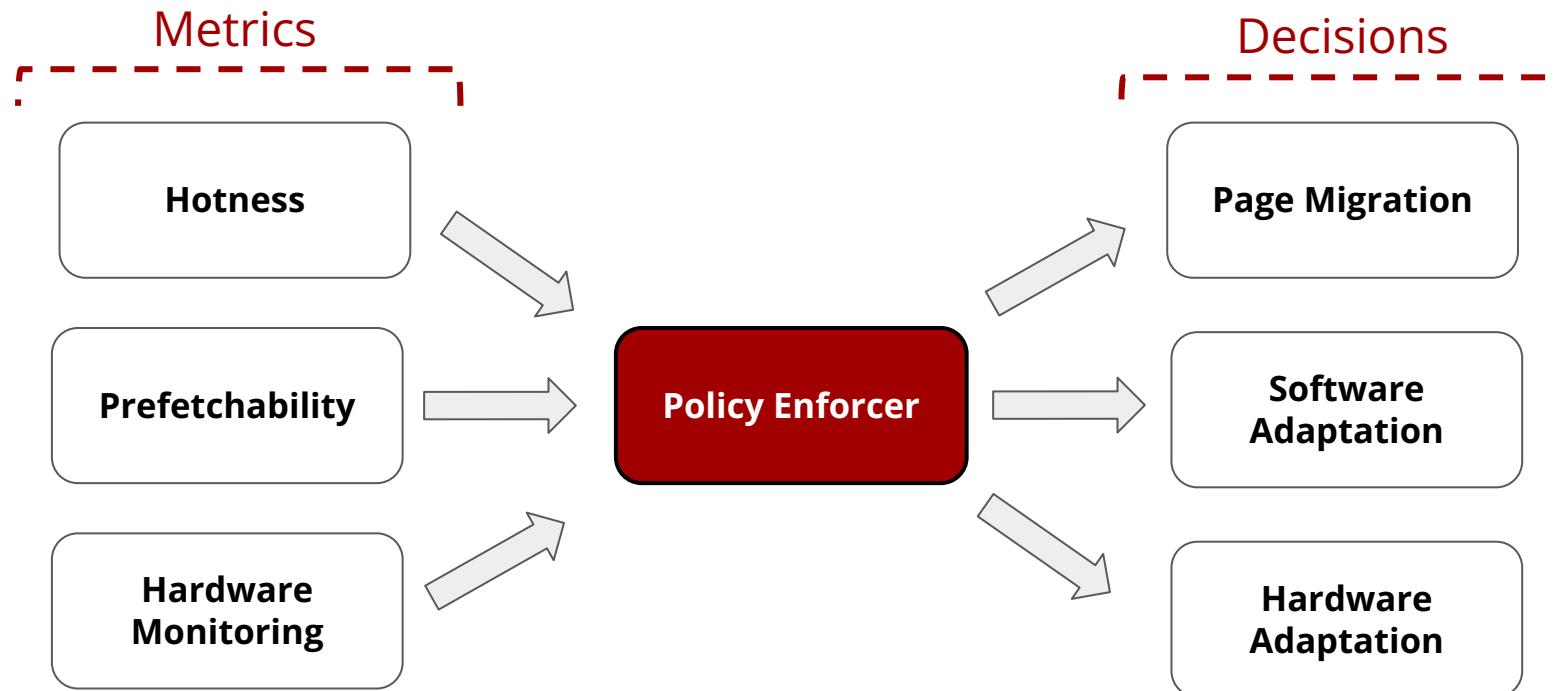
Observe dynamic behaviours

Adapts to the current state

Linden compiler

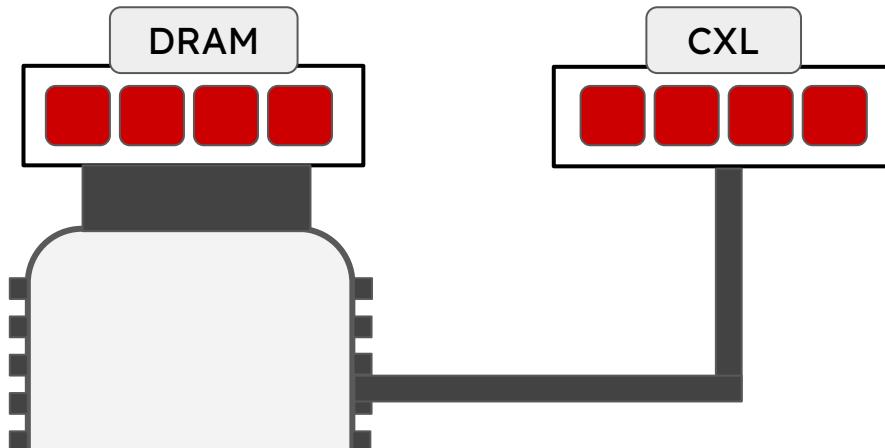


Linden's runtime responsibilities



Page migration

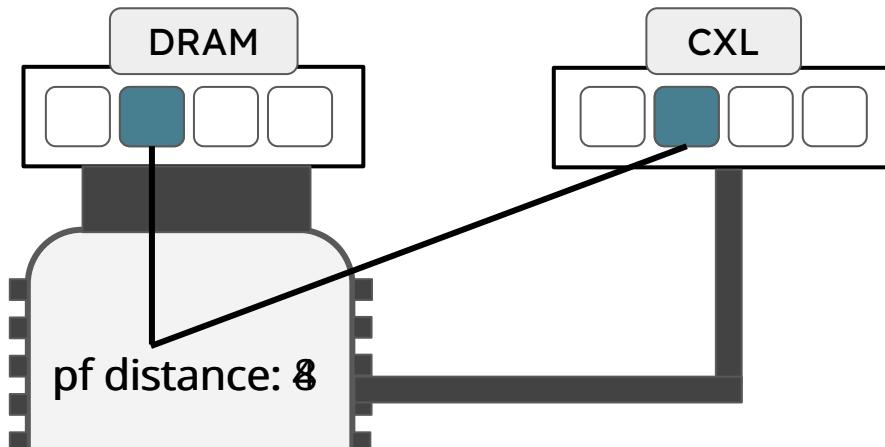
What happens if all regions are “hot” ?



- *Hotness + Prefetchability*
- Find memory regions that can **tolerate** the latency.
- Then migrate these pages to the CXL.

Software adaptation

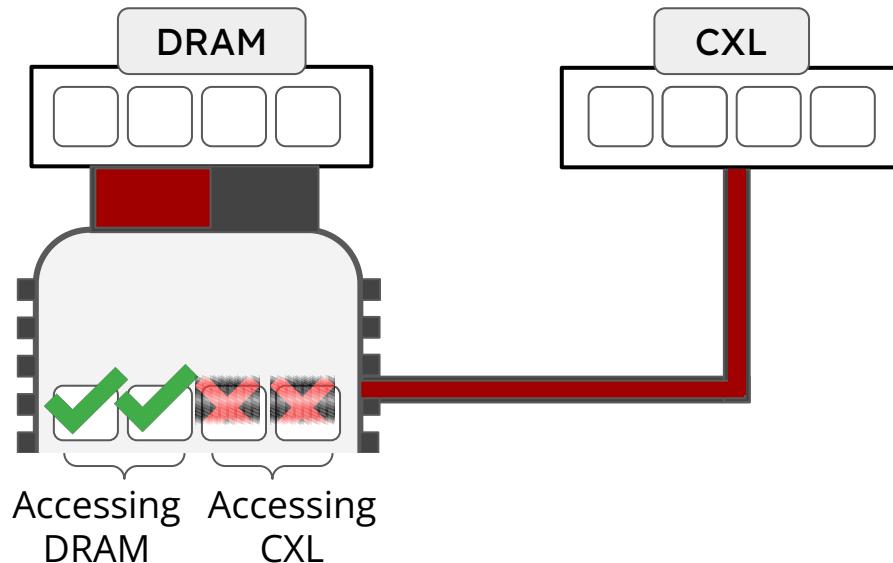
What if a region with prefetch migrated?



- Prefetch distance should be updated!
- Dynamically adjust the prefetch distances with considering the tier information.

Hardware adaptation

What happens under high contention.



- Disable hardware prefetching!
- Check which threads are accessing the CXL memory.

Linden: Tolerate It if You Cannot Reduce It

- Two techniques for tiered memory management:
 - Reduce understands the hotness.
 - Tolerate understands the prefetchability.
- Why toleration matters ?
 - Can help to hide CXL's latency
 - Careful monitoring is essential
 - Hardware and software adaptations needed

Thanks